

Stats Final Project

Name: Kanishk Yadav

UTEP Id: 80772034

Project Topic: Breast Cancer Classification

Objective: During my machine learning project, our main objective was to classify the tumour as malignant and benign using different machine learning techniques.

Contents:-

Data importing and Exploration.....2 - 6

The data that I used is provided by the UCI Machine Learning Repository and is also available as a callable data in the “sklean.datasets” library. I then performed data exploration meaning seeing the different data components by describing them and other methods.

Data visualization.....6 - 9

For data visualization, I used the machine learning libraries such as Seaborn and Matplotlib. We visualized data with graphs such as heatmap and scatterplot.

Model evaluation.....9 - 17

Here I start applying the actual machine learning algorithms to the model after I have explored and visualized the data. For my project, I used Support Vector Machine (SVM), Random Forest and Logistic Regression.

Conclusion.....17 - 18

I applied the 3 machine learning techniques and their results are available in the later parts of the document.

Appendix.....18 - 19

Important Codes

Data importing and Exploration

>>> I imported the data using from the sklearn preloaded datasets and the imported data was stored in the variable "cancer".

The "cancer" was then described for exploring and I got the following output: -

```
.. _breast_cancer_dataset:

Breast cancer wisconsin (diagnostic) dataset
-----

**Data Set Characteristics:**

    :Number of Instances: 569

    :Number of Attributes: 30 numeric, predictive attributes and the
    class

    :Attribute Information:
        - radius (mean of distances from center to points on the
        perimeter)
        - texture (standard deviation of gray-scale values)
        - perimeter
        - area
        - smoothness (local variation in radius lengths)
        - compactness (perimeter^2 / area - 1.0)
        - concavity (severity of concave portions of the contour)
        - concave points (number of concave portions of the contour)
        - symmetry
        - fractal dimension ("coastline approximation" - 1)

    The mean, standard error, and "worst" or largest (mean of
    the three
    worst/largest values) of these features were computed for
    each image,
    resulting in 30 features. For instance, field 0 is Mean
    Radius, field
    10 is Radius SE, field 20 is Worst Radius.

    - class:
        - WDBC-Malignant
        - WDBC-Benign

:Summary Statistics:

=====
radius (mean):      6.981  28.11
texture (mean):     9.71   39.28
perimeter (mean):   43.79  188.5
area (mean):        143.5  2501.0
=====
```

| | | |
|-------------------------------------|-------|--------|
| smoothness (mean): | 0.053 | 0.163 |
| compactness (mean): | 0.019 | 0.345 |
| concavity (mean): | 0.0 | 0.427 |
| concave points (mean): | 0.0 | 0.201 |
| symmetry (mean): | 0.106 | 0.304 |
| fractal dimension (mean): | 0.05 | 0.097 |
| radius (standard error): | 0.112 | 2.873 |
| texture (standard error): | 0.36 | 4.885 |
| perimeter (standard error): | 0.757 | 21.98 |
| area (standard error): | 6.802 | 542.2 |
| smoothness (standard error): | 0.002 | 0.031 |
| compactness (standard error): | 0.002 | 0.135 |
| concavity (standard error): | 0.0 | 0.396 |
| concave points (standard error): | 0.0 | 0.053 |
| symmetry (standard error): | 0.008 | 0.079 |
| fractal dimension (standard error): | 0.001 | 0.03 |
| radius (worst): | 7.93 | 36.04 |
| texture (worst): | 12.02 | 49.54 |
| perimeter (worst): | 50.41 | 251.2 |
| area (worst): | 185.2 | 4254.0 |
| smoothness (worst): | 0.071 | 0.223 |
| compactness (worst): | 0.027 | 1.058 |
| concavity (worst): | 0.0 | 1.252 |
| concave points (worst): | 0.0 | 0.291 |
| symmetry (worst): | 0.156 | 0.664 |
| fractal dimension (worst): | 0.055 | 0.208 |

=====

:Missing Attribute Values: None

:Class Distribution: 212 - Malignant, 357 - Benign

:Creator: Dr. William H. Wolberg, W. Nick Street, Olvi L. Mangasarian

:Donor: Nick Street

:Date: November, 1995

This is a copy of UCI ML Breast Cancer Wisconsin (Diagnostic) datasets.

<https://goo.gl/U2Uwz2>

Features are computed from a digitized image of a fine needle aspirate (FNA) of a breast mass. They describe characteristics of the cell nuclei present in the image.

Separating plane described above was obtained using Multisurface Method-Tree (MSM-T) [K. P. Bennett, "Decision Tree Construction Via Linear Programming." Proceedings of the 4th Midwest Artificial Intelligence and Cognitive Science Society, pp. 97-101, 1992], a classification method which uses linear programming to construct a decision tree. Relevant features were selected using an exhaustive search in the space of 1-4 features and 1-3 separating planes.

The actual linear program used to obtain the separating plane in the 3-dimensional space is that described in:

[K. P. Bennett and O. L. Mangasarian: "Robust Linear Programming Discrimination of Two Linearly Inseparable Sets", Optimization Methods and Software 1, 1992, 23-34].

This database is also available through the UW CS ftp server:

```
ftp ftp.cs.wisc.edu
cd math-prog/cpo-dataset/machine-learn/WDBC/
```

.. topic:: References

- W.N. Street, W.H. Wolberg and O.L. Mangasarian. Nuclear feature extraction for breast tumor diagnosis. IS&T/SPIE 1993 International Symposium on Electronic Imaging: Science and Technology, volume 1905, pages 861-870, San Jose, CA, 1993.
- O.L. Mangasarian, W.N. Street and W.H. Wolberg. Breast cancer diagnosis and prognosis via linear programming. Operations Research, 43(4), pages 570-577, July-August 1995.
- W.H. Wolberg, W.N. Street, and O.L. Mangasarian. Machine learning techniques to diagnose breast cancer from fine-needle aspirates. Cancer Letters 77 (1994) 163-171.

The above output tells me that there are 569 instances. Out of these 569, 212 are malignant and 357 are benign tumours. Since our data didn't have any missing attributed, we didn't need to replace or remove any NaN's or drop any rows.

There are in total 30 attributes, or we can say feature in the data. Some of them have been listed in the summary statistic table in the above output.

Other information such as References, etc., has also been provided.

>>> Other operations for exploration of data were further performed which were performed to see things such as "target names", "feature names" and seeing the shape of the data which was (569, 30)

>>> Further I converted the "cancer" data stored variable into a data frame "df_cancer". The ".head()" command for df_cancer gave following output

```
[ ] df_cancer.head()
```

| | mean radius | mean texture | mean perimeter | mean area | mean smoothness | mean compactness | mean concavity | mean concave points | mean symmetry | mean fractal dimension | ... |
|---|----------------|-----------------|-------------------|--------------|--------------------|---------------------|-------------------|---------------------------|------------------|------------------------------|-----|
| 0 | 17.99 | 10.38 | 122.80 | 1001.0 | 0.11840 | 0.27760 | 0.3001 | 0.14710 | 0.2419 | 0.07871 | ... |
| 1 | 20.57 | 17.77 | 132.90 | 1326.0 | 0.08474 | 0.07864 | 0.0869 | 0.07017 | 0.1812 | 0.05667 | ... |
| 2 | 19.69 | 21.25 | 130.00 | 1203.0 | 0.10960 | 0.15990 | 0.1974 | 0.12790 | 0.2069 | 0.05999 | ... |
| 3 | 11.42 | 20.38 | 77.58 | 386.1 | 0.14250 | 0.28390 | 0.2414 | 0.10520 | 0.2597 | 0.09744 | ... |
| 4 | 20.29 | 14.34 | 135.10 | 1297.0 | 0.10030 | 0.13280 | 0.1980 | 0.10430 | 0.1809 | 0.05883 | ... |

5 rows × 31 columns

| ... | worst texture | worst perimeter | worst area | worst smoothness | worst compactness | worst concavity | worst concave points | worst symmetry | worst fractal dimension | target |
|-----|------------------|--------------------|---------------|---------------------|----------------------|--------------------|----------------------------|-------------------|-------------------------------|--------|
| ... | 17.33 | 184.60 | 2019.0 | 0.1622 | 0.6656 | 0.7119 | 0.2654 | 0.4601 | 0.11890 | 0.0 |
| ... | 23.41 | 158.80 | 1956.0 | 0.1238 | 0.1866 | 0.2416 | 0.1860 | 0.2750 | 0.08902 | 0.0 |
| ... | 25.53 | 152.50 | 1709.0 | 0.1444 | 0.4245 | 0.4504 | 0.2430 | 0.3613 | 0.08758 | 0.0 |
| ... | 26.50 | 98.87 | 567.7 | 0.2098 | 0.8663 | 0.6869 | 0.2575 | 0.6638 | 0.17300 | 0.0 |
| ... | 16.67 | 152.20 | 1575.0 | 0.1374 | 0.2050 | 0.4000 | 0.1625 | 0.2364 | 0.07678 | 0.0 |

Here we see that the our “df_cancer” dataframe has dimension of 5 X 31 and also we observed that the “target” variable has 0’s and 1’s in the column (i.e., 0 for malignant and 1 for benign). By default .head() function gives only top 5 rows.

Data Visulization


>>> During this portion of my project, I visualized the data to see how all the variables compare to each other and how machine learning is working. We used libraries seaborn and matplotlib.

We created a pairplot using seaborn to where hue was set to “target” variable. (Hue: Variable in data to map plot aspects to different colors and the variables were set to be: 'mean radius', 'mean texture', 'mean area', 'mean perimeter', 'mean smoothness'. The output for the pairplot is as follows:-




Here we can see the hue set by us “target”, gave the colours to our targets 0 (Malignant) as Blue and 1 (Benign) as Orange. We can see the interpretations here such as that malignant tumour seem to have higher mean perimeter and mean radius then as compared to benign, however, mean smoothness of benign variable is more that of malignant.

>>> We then drew the count plot to check the count of malignant and benign tumour. We got the following output for it:-



It is observable here that the benign tumour (orange) has higher count (which was also seen when we described that the data).

>>> We then drew a correlation matrix for df_cancer dataframe using seaborn library's heatmap, we got the following output:-



Here it can be observed that the heatmap is showing the relation between all the variables. We can see that lighter the colour of variable, the more correlation they have. The diagonal line in the centre with all the 1's depict the correlation among the same features.

Evaluating the model

>>> For evaluating the model, first we had to split the data into training and test set. Here we split the data as follows:-

X is assigned as the part of the “df_cancer” such that it take all the features except the “target” and “y” variable takes all the target variable.

```
X = df_cancer.drop(['target'],axis=1)
X
```

| | mean radius | mean texture | mean perimeter | mean area | mean smoothness | mean compactness | mean concavity | mean concave points | mean symmetry | mean fractal dimension | ... | worst radius | worst texture | worst perimeter | worst area |
|-----|----------------|-----------------|-------------------|--------------|--------------------|---------------------|-------------------|---------------------------|------------------|------------------------------|-----|-----------------|------------------|--------------------|---------------|
| 0 | 17.99 | 10.38 | 122.80 | 1001.0 | 0.11840 | 0.27760 | 0.30010 | 0.14710 | 0.2419 | 0.07871 | ... | 25.380 | 17.33 | 184.60 | 2019.0 |
| 1 | 20.57 | 17.77 | 132.90 | 1326.0 | 0.08474 | 0.07864 | 0.08690 | 0.07017 | 0.1812 | 0.05667 | ... | 24.990 | 23.41 | 158.80 | 1956.0 |
| 2 | 19.69 | 21.25 | 130.00 | 1203.0 | 0.10960 | 0.15990 | 0.19740 | 0.12790 | 0.2069 | 0.05999 | ... | 23.570 | 25.53 | 152.50 | 1709.0 |
| 3 | 11.42 | 20.38 | 77.58 | 386.1 | 0.14250 | 0.28390 | 0.24140 | 0.10520 | 0.2597 | 0.09744 | ... | 14.910 | 26.50 | 98.87 | 567.7 |
| 4 | 20.29 | 14.34 | 135.10 | 1297.0 | 0.10030 | 0.13280 | 0.19800 | 0.10430 | 0.1809 | 0.05883 | ... | 22.540 | 16.67 | 152.20 | 1575.0 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 564 | 21.56 | 22.39 | 142.00 | 1479.0 | 0.11100 | 0.11590 | 0.24390 | 0.13890 | 0.1726 | 0.05623 | ... | 25.450 | 26.40 | 166.10 | 2027.0 |
| 565 | 20.13 | 28.25 | 131.20 | 1261.0 | 0.09780 | 0.10340 | 0.14400 | 0.09791 | 0.1752 | 0.05533 | ... | 23.690 | 38.25 | 155.00 | 1731.0 |
| 566 | 16.60 | 28.08 | 108.30 | 858.1 | 0.08455 | 0.10230 | 0.09251 | 0.05302 | 0.1590 | 0.05648 | ... | 18.980 | 34.12 | 126.70 | 1124.0 |
| 567 | 20.60 | 29.33 | 140.10 | 1265.0 | 0.11780 | 0.27700 | 0.35140 | 0.15200 | 0.2397 | 0.07016 | ... | 25.740 | 39.42 | 184.60 | 1821.0 |
| 568 | 7.76 | 24.54 | 47.92 | 181.0 | 0.05263 | 0.04362 | 0.00000 | 0.00000 | 0.1587 | 0.05884 | ... | 9.456 | 30.37 | 59.16 | 268.6 |

569 rows × 30 columns

```
y = df_cancer['target']
y
```

```
0    0.0
1    0.0
2    0.0
3    0.0
4    0.0
...
564  0.0
565  0.0
566  0.0
567  0.0
568  1.0
```

Name: target, Length: 569, dtype: float64

Using “train_test_split” from sklearn.model_selection,, we split the X and y as follows in the output:-

```
from sklearn.model_selection import train_test_split
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.20, random_state=5)
```


Model 1) Support Vector Machine:

Here we first stored the SVM algorithm into a variable and then applied to the `X_test` and `y_test`. Upon seeing the score of the SVC on `X_test` and `y_test`, we got the following output:-

```
svc_model.score(X_test, y_test)
```

```
0.9385964912280702
```

>>> Confusion matrix was then created to see the TP, TN, FP, FN. We created the confusion matrix using heatmap. Before that, we created a variable to store `X_test` called "`y_predict_SVM`". We got the following output:-




Here we observed that our model had very less amount of Type-I and Type-II error.

>>> We then further went on to make a classification report for the above model and got the following report for the classification report.

```
print(classification_report(y_test, y_predict_SVM))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0.0 | 1.00 | 0.85 | 0.92 | 48 |
| 1.0 | 0.90 | 1.00 | 0.95 | 66 |
| accuracy | | | 0.94 | 114 |
| macro avg | 0.95 | 0.93 | 0.94 | 114 |
| weighted avg | 0.94 | 0.94 | 0.94 | 114 |


>>> ROC Curve was made for the SVM:



>>> Further the model for the SVM was improved by using normalization:-

$$x' = \frac{x - \min(x)}{\max(x) - \min(x)}$$

Now, after normalizing the data (or we can say doing min-max scaling), we further made confusion matrix and classification report and got the following output with the scales "X_train_scaled":-




The heatmap above shows slight reduction in Type-I error after model improvement of SVM by normalization and improvement in true positives.

Further the classification report after the normalization gave slight improvements too with increase in accuracy from 94 to 98.

```
print(classification_report(y_test,y_predict_svm_1))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0.0 | 1.00 | 0.96 | 0.98 | 48 |
| 1.0 | 0.97 | 1.00 | 0.99 | 66 |
| accuracy | | | 0.98 | 114 |
| macro avg | 0.99 | 0.98 | 0.98 | 114 |
| weighted avg | 0.98 | 0.98 | 0.98 | 114 |

>>> ROC Curve was made for the SVM after Normalization:



Model 2) Random Forest:-

We then used an ensemble method called Random Forest.

Here again the, upon using the algorithm for random forest by storing it into a variable and applying it on the data, we got following score, confusion matrix and classification report. In our model we set the following parameters:


`n_estimators = 100`

`n_jobs = -1` (meaning the CPU is instructed to use as many cores as available for use)

`random_state = 100`

```
random_forest.score(X_test, y_test)
```

```
0.9736842105263158
```




The confusion matrix clearly shows the Type-I and Type-II errors.

```
print(classification_report(y_test, y_pred_randomforest))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0.0 | 0.98 | 0.96 | 0.97 | 48 |
| 1.0 | 0.97 | 0.98 | 0.98 | 66 |
| accuracy | | | 0.97 | 114 |
| macro avg | 0.97 | 0.97 | 0.97 | 114 |
| weighted avg | 0.97 | 0.97 | 0.97 | 114 |

We then got the classification report as above, we can see that the accuracy for it happens to be 97 which is very good.

>>> ROC Curve for the Random Forest:



Model 3) Logistic Regression


>>> We first imported the library for logistic regression and then stored the algorithm of logistic regression in a variable. We then fit the data (i.e., `X_train`, `y_train`) to the algorithm.

We got the following score for it:-

```
logmodel.score(X_test, y_test)
```

```
0.956140350877193
```

>>> Further we created the confusion matrix using the heat map to see the TP, TN, FP and FN. We observed that the Type – I and the Type – II errors were miniscule in count. True positives and False Negatives are very justifiable in it.




>>> We once again created classification matrix for the logistic regression to see accuracy of our model and we got the following scores:-

```
print(classification_report(y_test,y_pred_logistic_regression))
```


| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0.0 | 1.00 | 0.90 | 0.95 | 48 |
| 1.0 | 0.93 | 1.00 | 0.96 | 66 |
| accuracy | | | 0.96 | 114 |
| macro avg | 0.96 | 0.95 | 0.95 | 114 |
| weighted avg | 0.96 | 0.96 | 0.96 | 114 |

>>> ROC for the logistic regression:



Comparison of 3 models accuracies

After getting the accuracies for all the three models, we made a scatterplot of those for comparison of the 3 models we got:-



From the above scatterplot, it's clearly observable that the accuracy of SVM are highest followed by Random Forest and Logistic Regression.

Conclusion

From analysing the data and applying all the algorithms, we were able to find that following was the order of the models to classify the given data correctly into benign and malignant tumours.

Support Vector Machine (with accuracy of 0.98 after normalization) > Random Forest (accuracy of 0.97) > Logistic regression (accuracy of 0.96)

We can conclude that, since all the models had >0.90 accuracy, it's clear that our model was able to classify the data to very good extent. Also, the heatmaps for the confusion matrix showed that there so less Type – I and Type –II errors in all three of our models suggesting less false positives and false negatives.

We can also conclude an obvious fact that SVM and Random Forest performed better as they are complex algorithms as compared to Logistic Regression which is much simpler as compared to SVM and Random Forest.

Appendix

Important codes:-

>>> Making Dataframe:

```
df_cancer = pd.DataFrame(np.c_[cancer['data'], cancer['target']], columns = np.append(cancer['feature_names'], ['target']))
```

>>> Splitting data into training and testing:

```
from sklearn.model_selection import train_test_split
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.20, random_state=5)
```

>>> SVC, Random Forest and Logistic Regression importing and storing:

```
from sklearn.svm import SVC
```

```
from sklearn.metrics import classification_report, confusion_matrix
```

```
svc_model = SVC()
```

```
from sklearn.ensemble import RandomForestClassifier
```

```
random_forest = RandomForestClassifier(n_estimators=100, n_jobs = -
1, random_state = 100)
```

```
from sklearn.linear_model import LogisticRegression
logmodel = LogisticRegression()
```

>>> Normalization:

```
min_train = X_train.min()
```

```
min_train
```

```
range_train = (X_train - min_train).max()
range_train
```

```
X_train_scaled = (X_train - min_train)/range_train
```

```
X_train_scaled
```

>>> Correlation Matrix:-

```
plt.figure(figsize=(20,10))
```

```
sns.heatmap(df_cancer.corr(), annot=True)
```

>>> ACCURACY COMAPARISONS:-

```
import seaborn as sns
```

```
sns.scatterplot(x = (0.98, 0.97, 0.96), y = ("SVM", "Random Forest", "L
ogistic Regression"))
```